

The background is a dark blue gradient with glowing blue circuit lines and dots. On the left, there is a glowing purple and blue rectangular area containing the letters 'AI' in a large, white, sans-serif font. The 'A' has a bright purple glow, and the 'I' has a bright blue glow. The background also features a grid of small white dots and lines, resembling a circuit board or a data network.

# AI

# LOCAL SEARCH AND OPTIMIZATION PROBLEMS

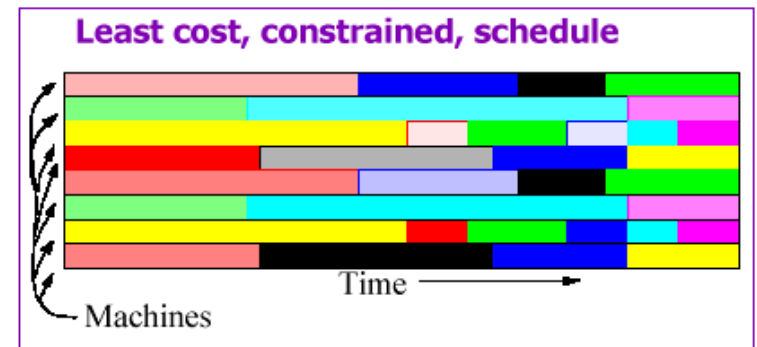
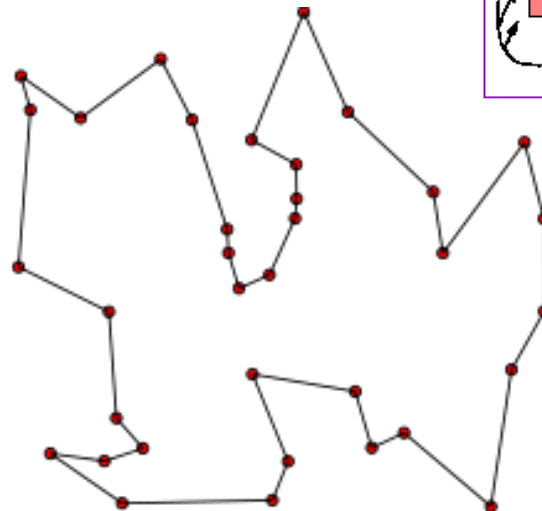
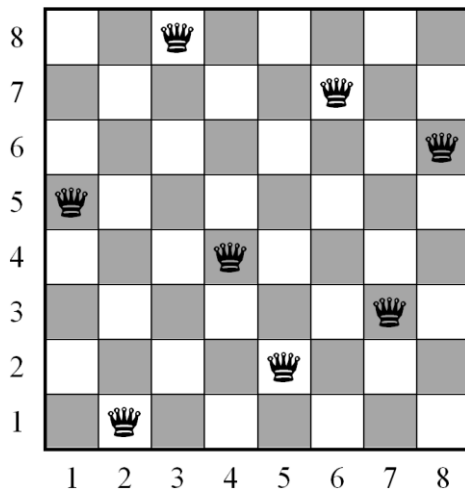
Nguyễn Ngọc Thảo – Nguyễn Hải Minh  
{nnthao, nhminh}@fit.hcmus.edu.vn

# Outline

- Local search and optimization problems
- Local search algorithms
- Evolutionary algorithms

# Paths are not always required

- There are applications that **only the final state matters**, not the path to get there.
  - Integrated-circuit design, factory floor layout, job shop scheduling, etc.

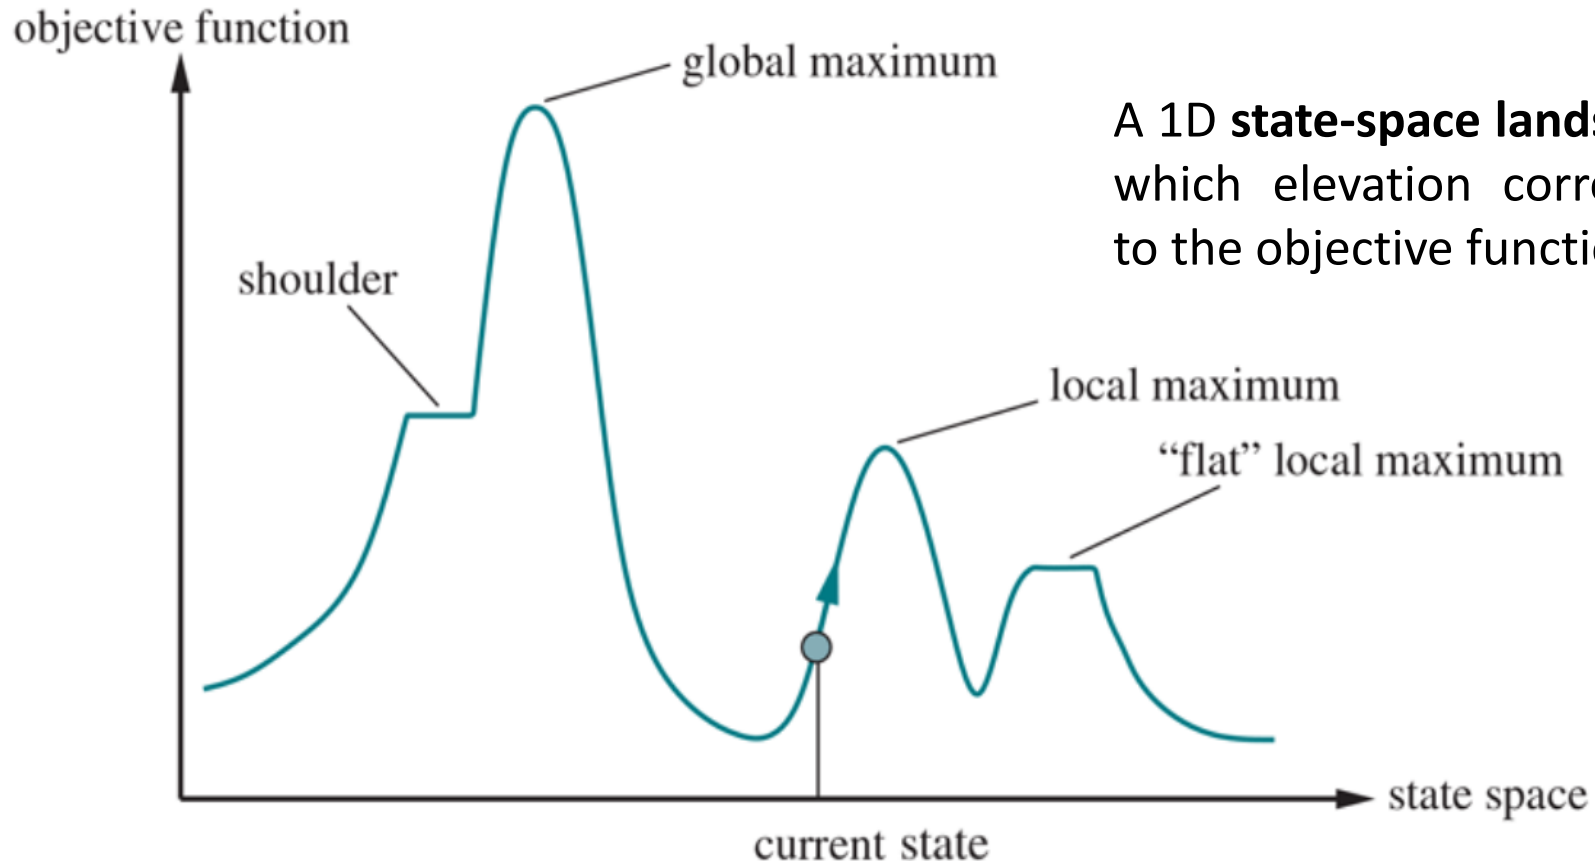


# Local search algorithms

- These algorithms **navigate from a start state to neighbors, without tracking the paths, nor the set of reached states.**
- **Not systematic**
  - They might never explore a portion of the search space where a solution resides.
- Local search has two key advantages:
  - Use very **little memory**
  - Find **reasonable solutions in large or infinite state spaces** for which systematic algorithms are unsuitable.

# Local search algorithms

- Local search can solve **optimization problems**, which finds the **best state according to an objective function**.



# Hill climbing search

- The algorithm **heads in the direction that gives the steepest ascent** and **terminates when it reaches a “peak”**.
  - Peak = a state where no neighbor has a higher value.

```
function HILL-CLIMBING(problem) returns a state that is a local maximum  
  current  $\leftarrow$  problem.INITIAL  
  while true do  
    neighbor  $\leftarrow$  a highest-valued successor of current  
    if VALUE(neighbor)  $\leq$  VALUE(current) then return current  
    current  $\leftarrow$  neighbor
```

Hill-climbing tracks only one current state and, on each iteration, moves to the neighboring state with highest value.

# Hill climbing for 8-queens problem

- **Complete-state formulation:** all queens on the board, one per column
- **Successor function:** move a queen to another square in the same column  $\rightarrow$  each state has  $8 \times 7 = 56$  successors
- $h(n)$  = the number of pairs of queens that are attacking each other  $\rightarrow$  the global minimum has  $h(n)=0$

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♔	13	16	13	16
♔	14	17	15	♔	14	16	16
17	♔	16	18	15	♔	15	♔
18	14	♔	15	15	14	♔	16
14	14	13	17	12	14	12	18

The board shows the value of for each possible successor obtained by moving a queen within its column.

The current state  $n$  has  $h(n) = 17$   
 $(c_1 c_2 c_3 c_4 c_5 c_6 c_7 c_8) = (4\ 3\ 2\ 5\ 4\ 3\ 2\ 3)$


There are 8 moves that are tied for best, with  $h = 12$ . Hill climbing algorithm will pick one of these.

# Hill climbing: Suboptimality

- Hill climbing is also called **greedy local search**.
  - It grabs a good neighbor state without thinking where to go next.
- It can easily improve a bad state and hence make rapid progress toward a solution.
  - 8-queens: 17 million states, 4 steps on average when succeeds and 3 when get stuck.
- Hill climbing can get stuck in **local maxima**, **ridges**, or **plateaus**




# Local extrema and ridges



- Current state (1 6 2 5 7 4 8 3) has  $h(n) = 1$
- Every successor has a higher cost  
→ local minimum

The grid of states (dark circles) is laid on a ridge rising from left to right, creating a sequence of local maxima

From each local maximum, all the available actions point downhill.



# Local extrema and ridges

- Real-world problem or NP-hard problems typically have an exponential number of local maxima to get stuck on.

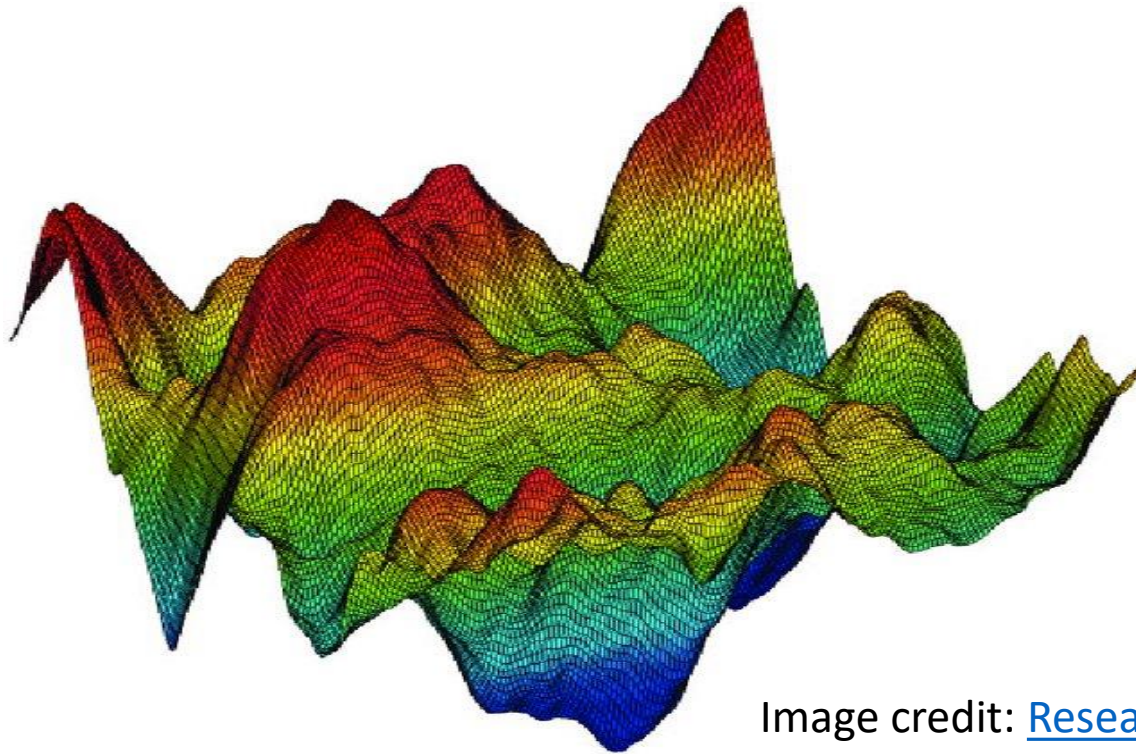


Image credit: [Research Gate](#)

# Overcome the suboptimality


- A **sideways move** lets the agent keep going in the plateau.
  - It does not work on a flat local maximum.
- The number of consecutive sideways moves should be limited to avoid non-stop wander.
- This approach raises the percentage of problem instances solved by hill climbing.
  - E.g., for 8-queens problem: from 14% to 94%.
  - **Success comes at a cost:** roughly 21 steps for each successful instance and 64 for each failure

# Overcome the suboptimality

- **Stochastic hill climbing** chooses at random from among the uphill moves.
  - The probability of selection can vary with the steepness of the move.
  - Slower convergence, yet better solutions in some state landscapes
- **First-choice hill climbing** generates successors randomly until obtaining one that is better than the current state.
  - Suitable when a state has many (e.g., thousands) of successors
- **Random-restart hill climbing** conducts several hill-climbing searches from random initial states, until a goal is found.
  - If each search has a probability  $p$  of success, the expected number of restarts required is  $1/p$ .

# Quiz 01: 4-queens problem


- Consider the following 4-queens problem



- Apply hill-climbing to find a solution, using the heuristic “*The number of pairs of queens attacking each other.*”

# Simulated annealing

- Combine **hill climbing** with a **random walk** in some way that yields both **efficiency** and **completeness**



Shake hard (i.e., at a high temperature) and then gradually reduce the intensity of shaking (i.e., lower the temperature)

# Simulated annealing

a mapping function from  
time to “temperature”

**function** SIMULATED-ANNEALING(*problem*, *schedule*) **returns** a solution state

*current*  $\leftarrow$  *problem*.INITIAL

**for**  $t = 1$  **to**  $\infty$  **do**

$T \leftarrow \text{schedule}(t)$


**if**  $T = 0$  **then return** *current*

*next*  $\leftarrow$  a randomly selected successor of *current*

$\Delta E \leftarrow \text{VALUE}(\text{next}) - \text{VALUE}(\text{current})$

**if**  $\Delta E > 0$  **then** *current*  $\leftarrow$  *next*

**else** *current*  $\leftarrow$  *next* only with probability  $e^{\Delta E/T}$




# Local beam search

- Keeping just one node in memory might seem to be an extreme reaction to the problem of memory limitations.
- The algorithm **keeps track of  $k$  states** rather than just one.
- It begins with randomly generated states.
- At each step, all the successors of all  $k$  states are generated
- If any one is a goal, the algorithm halts. Otherwise, it selects the best successors from the complete list and repeats.



# Local beam search




- The algorithm quickly abandons unfruitful searches and moves its resources to where the most progress is being made.

# Local beam search

- Useful information is passed among the parallel search threads → major difference from random-restart search
- The algorithm possibly suffers from a lack of diversity among the  $k$  states.
  - The states can become clustered in a small region of the state space → an expensive version of hill climbing.
- Stochastic beam search can alleviate the above problem by randomly picking  $k$  successors following their values.

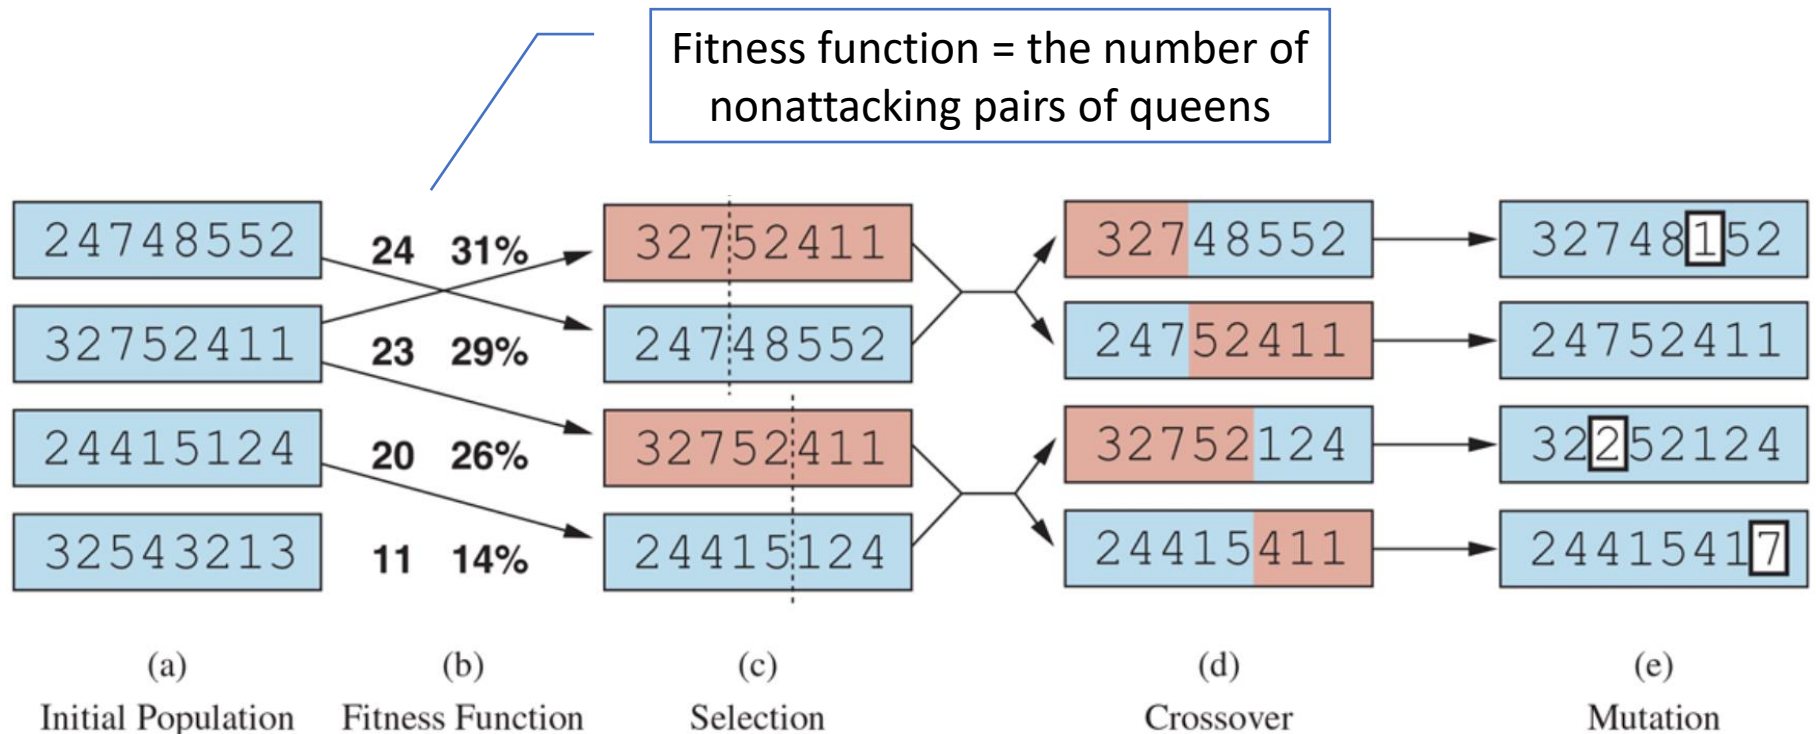
# Evolutionary algorithms

- Variants of stochastic beam search, explicitly motivated by the metaphor of natural selection in biology



There is a population of individuals (states). The fittest (highest value) individuals produce offspring (successor states) that populate the next generation.

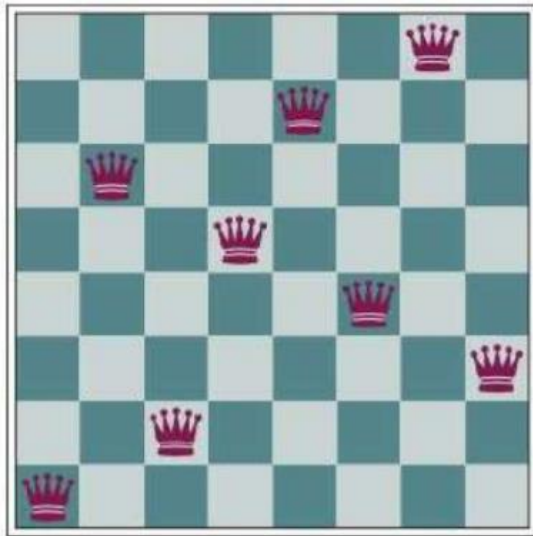
# Evolutionary algorithms



A genetic algorithm, illustrated for digit strings representing 8-queens states. The initial population in (a) is ranked by a fitness function in (b) resulting in pairs for mating in (c). They produce offspring in (d), which are subject to mutation in (e).

# Evolutionary algorithms

- There are endless forms of evolutionary algorithms.
- **The representation of an individual**
  - **Genetic algorithms:** a string over a finite alphabet.
  - Genetic programming: a computer program. Evolution strategies: a sequence of real numbers.



- Digit representation  
 $\langle 1\ 6\ 2\ 5\ 7\ 4\ 8\ 3 \rangle$
- Binary representation  
 $\langle 000\ 101\ 001\ 100\ 110\ 011\ 111\ 010 \rangle$

# Evolutionary algorithms

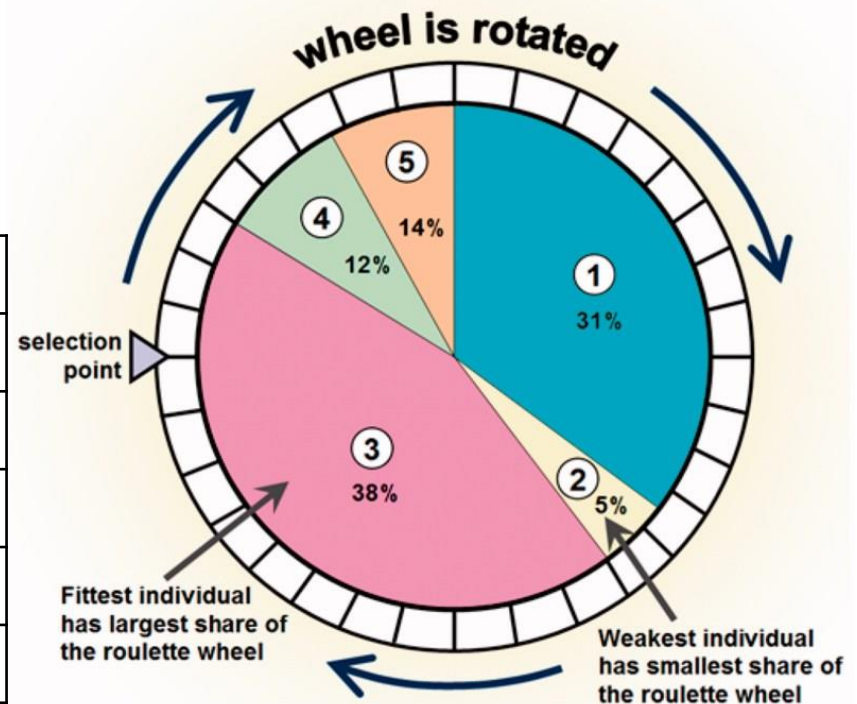
- The size of the population
  - A population is a set of  $k$  randomly generated states to begin with.
- **Fitness function:** an objective function that rates each state
  - The higher values, the better states
- The mixing number  $\rho$ : number of parents that come together to form offsprings
  - $\rho = 1$ : stochastic beam search.  $\rho = 2$ : most common.

# Evolutionary algorithms

- The selection process: choose individuals as parents of the next generation
  - Individuals can be chosen with probabilities proportional to their fitness score.

## The Roulette wheel method

Individual	Fitness percentage (%)
1	31
2	5
3	38
4	12
5	14



# Evolutionary algorithms

- The recombination procedure to form children
  - It randomly selects a **crossover point** to split each of the parent strings and recombines the parts to form two children.

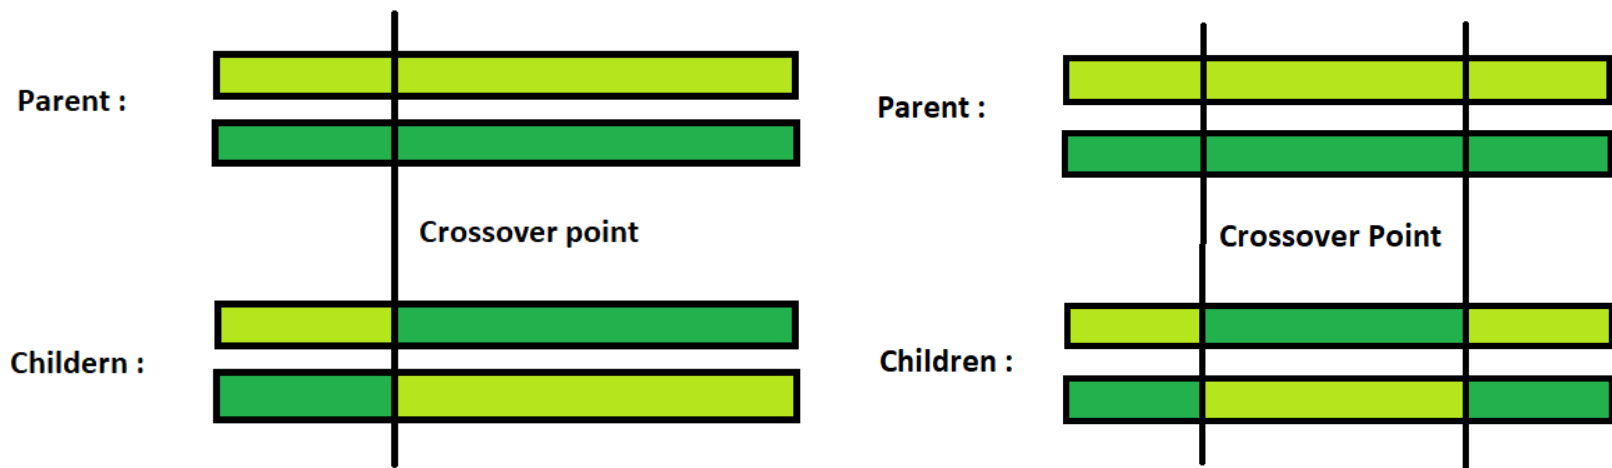
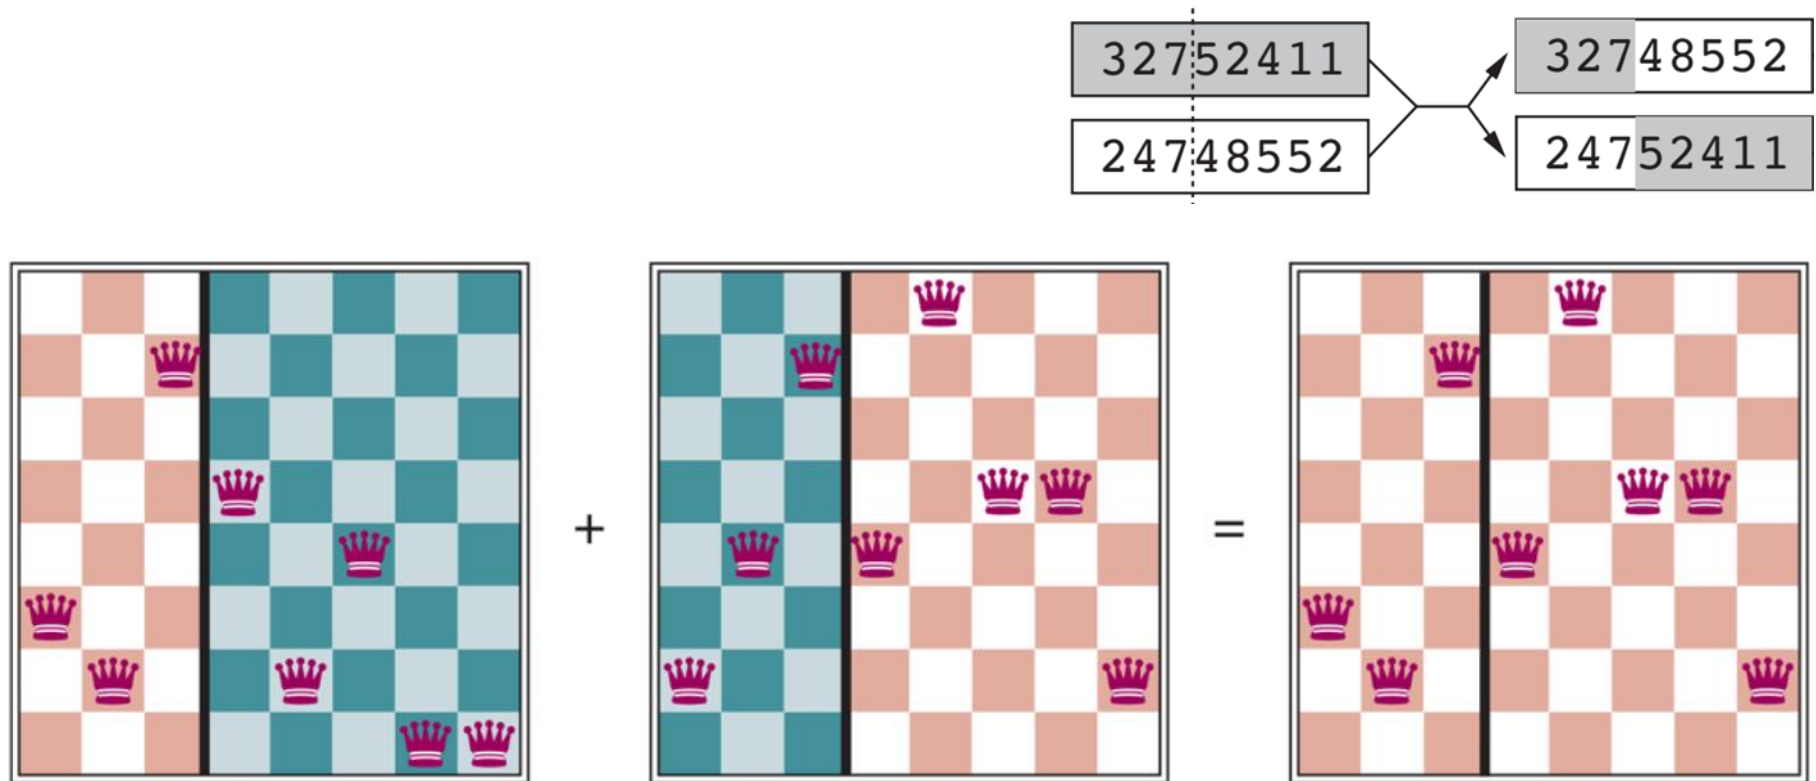


Image credit: [Geeksforgeeks](https://www.geeksforgeeks.org/)

- Crossover frequently takes large steps in the state space early in the search process.



# Recombination step: An example



The 8-queens states corresponding to the two parents (left and middle) and the first offspring (right). The green columns are lost in the crossover step and the red columns are retained.

# Evolutionary algorithms

- **The mutation rate:** determine how often offspring have random mutations to their representation.
  - Every bit in the individual's composition is flipped with probability equal to the mutation rate.
- **The makeup of the next generation:** only the newly formed offspring, or a few top-scoring parents also included.
  - **Elitism practice:** guarantee that overall fitness will never decrease over time
  - **Culling practice:** All individuals below a given threshold are discarded, which can lead to a speedup (Baum et al., 1995).

# Quiz 02: Calculate fitness scores

- Consider the 4-queens problem, in which each state has 4 queens, one per column, on the board. The state can be represented in genetic algorithm as a sequence of 4 digits, each of which denotes the position of a queen in its own column (from 1 to 4).



- $Fit(n)$  = the number of non-attacking pairs of queens*
- Let the current generation includes 4 states:  
$$S1 = 2341; \quad S2 = 2132; \quad S3 = 1232; \quad S4 = 4321.$$
- Calculate the value of  *$Fit(n)$*  for the given states and the probability that each of them will be chosen in the “selection” step.

```
function GENETIC-ALGORITHM(population, fitness) returns an individual
  repeat
    weights  $\leftarrow$  WEIGHTED-BY(population, fitness)
    population2  $\leftarrow$  empty set
    for i = 1 to SIZE(population) do
      parent1, parent2  $\leftarrow$  WEIGHTED-RANDOM-CHOICES(population, weights, 2)
      child  $\leftarrow$  REPRODUCE(parent1, parent2)
      if (small random probability) then child  $\leftarrow$  MUTATE(child)
      add child to population2
    population  $\leftarrow$  population2
  until some individual is fit enough, or enough time has elapsed
  return the best individual in population, according to fitness
```

```
function REPRODUCE(parent1, parent2) returns an individual
  n  $\leftarrow$  LENGTH(parent1)
  c  $\leftarrow$  random number from 1 to n
  return APPEND(SUBSTRING(parent1, 1, c), SUBSTRING(parent2, c+1, n)
```

# An evaluation of Genetic algorithms

- Crossover gives better random exploration than local search.
- Rely on very little domain knowledge
- Large number of “tunable” parameters
  - Difficult to replicate performance from one problem to another
- Lack of good empirical studies comparing to simpler methods
  - Useful on some (small?) sets of problem, yet no convincing evidence that GAs are better than hill-climbing w/random restarts in general.

